

Exercises for Introduction to Python

Prof. Dr. Markus Löcher

06/27/2023

Table of contents

Preface	4
1 Lists, Loops, Conditions	5
1.0.1 Manipulating lists of lists	5
1.0.2 Automation by iterating	6
1.0.3 Conditions	6
2 Functions, Dictionaries	7
2.0.1 Functions	7
2.0.2 Dictionaries	11
2.0.3 Introduction to numpy	13
3 Numpy Arrays, Randomness	15
3.0.1 Numpy Arrays	15
3.0.2 Random Data Generation	15
3.0.3 Simulating Probabilistic Events	16
4 Probabilistic Events	17
4.0.1 Simulating Probabilistic Events	17
5 Pandas	18
5.0.1 Gapminder	19
6 Missing Data	20
6.0.1 Tasks:	23
7 Regression, seaborn	24
7.1 Seaborn Graphs	25
7.2 Violin Plots	26
7.2.1 Histograms	27
7.2.2 Tasks:	29
7.3 Extra Credit	30
8 String Exercises	32
8.1 Part I	32
8.1.1 NOTE!	32

8.2	Part II (Netflix data)	32
8.3	Netflix	35
9	Subplots and Pivoting	36
9.1	Smoking Data	37
9.2	Netflix Data	39

Preface

This is a collection of exercises that accompany the python course.

1 Lists, Loops, Conditions

In this lab we will get to know and become experts in:

1. Lists
 - DataCamp, [Introduction to Python](#), Chap 2
2. Loops
 - DataCamp, [Intermediate Python](#), Chap 4
3. Conditions
 - DataCamp, [Intermediate Python](#), Chap 3

1.0.1 Manipulating lists of lists

The following list of lists contains names of sections in a house and their area.

1. Extract the area corresponding to kitchen
2. String Tasks:
 - Extract the first letters of each string
 - Capitalize all strings
 - Replace all occurrences of “room” with “rm”
 - count the number of “l” in “hallway”
3. Insert a “home office” with area 10.75 after living room
4. Append the total area to the end of the list
5. **Boolean** operations:
 - Generate one True and one False by comparing areas
 - Generate one True and one False by comparing names

```
house = [['hallway', 11.25],  
         ['kitchen', 18.0],  
         ['living room', 20.0],  
         ['bedroom', 10.75],  
         ['bathroom', 9.5]]
```

1.0.2 Automation by iterating

for loops are a powerful way of automating MANY otherwise tedious tasks that repeat.

1. Repeat the tasks 2 and 4 from above by using a for loop
 - using `enumerate`
 - using `range`
2. Create two separate new lists which contain only the names and areas separately
3. [Clever Carl](#): Compute

$$\sum_{i=1}^{100} i$$

```
list(range(5))
```

```
[0, 1, 2, 3, 4]
```

1.0.3 Conditions

1. Find the **max** of the areas by using `if` inside a for loop
2. Print those elements of the list with
 - `area > 15`
 - strings that contain “room” (or “rm” after your substitution)

2 Functions, Dictionaries

In this lab we will get to know and become experts in:

1. [Functions](#)
 - DataCamp, [Introduction to Python](#), Chap 3
2. [Dictionaries](#)
 - DataCamp, [Intermediate Python](#), Chap 2
3. [Introduction to numpy](#)
 - DataCamp, [Introduction to Python](#), Chap 4

2.0.1 Functions

[Functions](#) are essential building blocks to **reuse code** and to **modularize code**.

We have already seen and used many **built-in functions/methods** such as `print()`, `len()`, `max()`, `round()`, `index()`, `capitalize()`, etc..

```
areas = [11.25, 18.0, 20.0, 10.75, 10.75, 9.5]
print(max(areas))
print(len(areas))
print(round(10.75,1))
print(areas.index(18.0))
```

```
20.0
6
10.8
1
```

But of course we want to define our own functions as well ! As a rule of thumb, if you anticipate needing to repeat the same or very similar code more than once, it may be worth writing a reusable function. Functions can also help make your code more readable by giving a name to a group of Python statements.

For example, we computed the BMI previously as follows:

```
height = 1.79
weight = 68.7
bmi = weight/height**2
print(bmi)
```

21.44127836209856

Functions are declared with the `def` keyword. A function contains a block of code with an optional use of the `return` keyword:

```
def compute_bmi(height, weight):
    return weight/height**2

compute_bmi(1.79, 68.7)
```

21.44127836209856

Each function can have *positional* arguments and *keyword* arguments. Keyword arguments are most commonly used to specify default values or optional arguments. For example:

```
def compute_bmi(height, weight, ndigits=2):
    return round(weight/height**2, ndigits)

print(compute_bmi(1.79, 68.7))
print(compute_bmi(1.79, 68.7, 4))
```

21.44

21.4413

2.0.1.1 Multiple Return Values

are easily possible in python:

```
def compute_bmi(height, weight, ndigits=2):
    bmi = round(weight/height**2, ndigits)
    #https://www.cdc.gov/healthyweight/assessing/index.html#:~:text=If%20your%20BMI%20is%2
    if bmi < 18.5:
```



```

        status="underweight"
    elif bmi <= 24.9:
        status="healthy"
    elif bmi <= 29.9:
        status="underweight"
    elif bmi >= 30:#note that a simple else would suffice here!
        status="obese"
    return bmi, status

print(compute_bmi(1.79, 68.7))
print(compute_bmi(1.79, 55))

```

```

(21.44, 'healthy')
(17.17, 'underweight')

```

Recall from the previous lab how we

1. found the largest room,
2. computed the sum of integers from 1 to 100

```

#find the maximum area:
areas = [11.25, 18.0, 20.0, 10.75, 10.75, 9.5]
currentMax = areas[0] # initialize to the first area seen

for a in areas:
    if a > currentMax:
        currentMax = a

print("The max is:", currentMax)

```

The max is: 20.0

```

#Clever IDB students: Compute the sum from 1 to 100:
Total =0

for i in range(101):#strictly speaking we are adding the first 0
    Total = Total + i
    #Total += i

print(Total)

```

2.0.1.2 Tasks

Write your own function

1. to find the min and max of a list
2. to compute the Gauss sum with default values $m = 1, n = 100$

$$\sum_{i=m}^n i$$

2.0.1.3 Namespaces and Scope

Functions seem straightforward. But one of the more confusing aspects in the beginning is the concept that we can have **multiple instances** of the same variable!

Functions can access variables created inside the function as well as those outside the function in higher (or even global) scopes. An alternative and more descriptive name describing a variable scope in Python is a *namespace*. Any variables that are assigned within a function by default are assigned to the local namespace. The local namespace is created when the function is called and is immediately populated by the function's arguments. After the function is finished, the local namespace is destroyed.

Examples:

```
height = 1.79
weight = 68.7
bmi = weight/height**2
#print("height, weight, bmi OUTSIDE the function:",height, weight,bmi)

def compute_bmi(h, w):
    height = h
    weight = w
    bmi = round(weight/height**2,2)
    status="healthy"
    print("height, weight, bmi INSIDE the function:",height, weight,bmi)
    print("status:", status)
    return bmi

compute_bmi(1.55, 50)

print("height, weight, bmi OUTSIDE the function:",height, weight,bmi)
#print(status)
```

```
height, weight, bmi INSIDE the function: 1.55 50 20.81
status: healthy
height, weight, bmi OUTSIDE the function: 1.79 68.7 21.44127836209856
```

2.0.2 Dictionaries

A [dictionary](#) is basically a **lookup table**. It stores a collection of key-value pairs, where key and value are Python objects. Each key is associated with a value so that a value can be conveniently retrieved, inserted, modified, or deleted given a particular key.

The dictionary or `dict` may be the most important built-in Python data structure. In other programming languages, dictionaries are sometimes called *hash maps* or *associative arrays*.

```
#This was the house defined as a list of lists:
house = [['hallway', 11.25],
        ['kitchen', 18.0],
        ['living room', 20.0],
        ['bedroom', 10.75],
        ['bathroom', 9.5]]

#Remember all the disadvantages of accessing elements

#Better as a lookup table:
house = {'hallway': 11.25,
        'kitchen': 18.0,
        'living room': 20.0,
        'bedroom': 10.75,
        'bathroom': 9.5}
```

```
europe = {'spain':'madrid', 'france' : 'paris'}
print(europe["spain"])
print("france" in europe)
print("paris" in europe)#only checks the keys!
europe["germany"] = "berlin"
print(europe.keys())
print(europe.values())
```

```
madrid
True
False
dict_keys(['spain', 'france', 'germany'])
dict_values(['madrid', 'paris', 'berlin'])
```

If you need to iterate over both the keys and values, you can use the `items` method to iterate over the keys and values as 2-tuples:

```
#print(list(europe.items()))

for country, capital in europe.items():
    print(capital, "is the capital of", country)
```

```
madrid is the capital of spain
paris is the capital of france
berlin is the capital of germany
```

Note: You can use integers as keys as well. However -unlike in lists- one should not think of them as positional indices!

```
#Assume you have a basement:
house[0] = 21.5
house
```

```
{'hallway': 11.25,
 'kitchen': 18.0,
 'living room': 20.0,
 'bedroom': 10.75,
 'bathroom': 9.5,
 0: 21.5}
```

```
#And there is a difference between the string and the integer index!
house["0"] = 30.5
house
```

```
{'hallway': 11.25,
 'kitchen': 18.0,
 'living room': 20.0,
 'bedroom': 10.75,
 'bathroom': 9.5,
 0: 21.5}
```

Categorize a list of words by their first letters as a dictionary of lists:

```

words = ["apple", "bat", "bar", "atom", "book"]

by_letter = {}

for word in words:
    letter = word[0]
    if letter not in by_letter:
        by_letter[letter] = [word]
    else:
        by_letter[letter].append(word)

```

```
{'a': ['apple', 'atom'], 'b': ['bat', 'bar', 'book']}
```

2.0.2.1 Tasks

1. Find the maximum of the areas of the houses
2. Remove the two last entries.
3. Write a function named `word_count` that takes a string as input and returns a dictionary with each word in the string as a key and the number of times it appears as the value.

2.0.3 Introduction to numpy

NumPy, short for Numerical Python, is one of the most important foundational packages for numerical computing in Python.

1. Vectorized, fast mathematical operations.
2. Key features of NumPy is its N-dimensional array object, or `ndarray`

```

height = [1.79, 1.85, 1.95, 1.55]
weight = [70, 80, 85, 65]

#bmi = weight/height**2

import numpy as np

height = np.array([1.79, 1.85, 1.95, 1.55])
weight = np.array([70, 80, 85, 65])

bmi = weight/height**2
np.round(bmi,2)

```

```
array([21.84700852, 23.37472608, 22.35371466, 27.05515088])
```

3 Numpy Arrays, Randomness

In this lab we will get to know and become experts in:

1. [Numpy Arrays](#)
 - Slicing and Accessing
 - Properly using axis
2. [Random Data Generation](#)
 - Random integers, permutations and sampling
3. [Simulating Probabilistic Events](#)
 -

3.0.1 Numpy Arrays

3.0.1.1 Tasks

1. Generate a sequence from 1 to 64
2. Print every other element
3. Using Boolean indexing: print only those numbers that are greater than 10
4. Reshape into a 8x8 matrix and print its “shape”
5. Compute the column and row sums

3.0.2 Random Data Generation

1. “Flip a fair coin” 20 times and save into an array. Note that instead of using “heads/tails” you should “code” the outcome as 0/1.
2. Randomly “draw” 2 integers without replacement from the sequence 1-5. Repeat this process 30 times and store the results in an array.
3. Compute the counts

3.0.3 Simulating Probabilistic Events

1. **Overbooking flights:** airlines
2. **Home Office** days: planning office capacities and minimizing social isolation

4 Probabilistic Events

More Simulations of Probabilistic Events

```
import numpy as np
from numpy.random import default_rng
```

4.0.1 Simulating Probabilistic Events

1. **Biased Coin:** Simulate 365 days with a $p = \frac{1}{4}$ chance of being sunny (=1). Hint: exploit the fact that p is a fraction!
2. **Birthday problem** Change the “birthday code” into a function with “n = number of people in a room” as an argument. (What other arguments might be useful?) Execute this function for $n = 10, 25, 50$.
3. **Overbooking flights:** Imagine an airline sold 105 tickets on a flight with 100 seats. Assuming there is a 10% no-show probability per passenger, “compute” (simulate) the probability that the airline will need to pay someone to not board.

5 Pandas

All about pandas

```
import numpy as np
import pandas as pd
!pip install gapminder
from gapminder import gapminder
```

```
height = np.array([1.79, 1.85, 1.95, 1.55])
weight = np.array([70, 80, 85, 65])
hw = np.array([height, weight]).transpose()
```

```
hw
```

```
array([[ 1.79, 70.  ],
       [ 1.85, 80.  ],
       [ 1.95, 85.  ],
       [ 1.55, 65.  ]])
```

```
df = pd.DataFrame(hw , columns = ["height", "weight"])
print(df)
```

```
   height  weight
0    1.79    70.0
1    1.85    80.0
2    1.95    85.0
3    1.55    65.0
```

```
df = pd.DataFrame(hw , columns = ["height", "weight"],
                  index = ["Peter", "Matilda", "Bee", "Bee"])
print(df)
```

	height	weight
Peter	1.79	70.0
Matilda	1.85	80.0
Bee	1.95	85.0
Bee	1.55	65.0

Can you extract:

0. All weights
1. Peter's height
2. Bee's full info
3. the average height
4. get all persons with height greater than 180cm

5.0.1 Gapminder

```
gapminder.head()
```

	country	continent	year	lifeExp	pop	gdpPercap
0	Afghanistan	Asia	1952	28.801	8425333	779.445314
1	Afghanistan	Asia	1957	30.332	9240934	820.853030
2	Afghanistan	Asia	1962	31.997	10267083	853.100710
3	Afghanistan	Asia	1967	34.020	11537966	836.197138
4	Afghanistan	Asia	1972	36.088	13079460	739.981106

Tasks

- find the unique years
- get all rows with year 1952
- get all rows from 1952:1962
- get all rows from Afghanistan to Albania

6 Missing Data

More pandas but this time on a real data set, namely the [kaggle Housing Data](#) which you can read directly from Google Drive

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
#!pip install gapminder
#from gapminder import gapminder
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
#read in the data
rootPath = "/content/drive/MyDrive/"#same for all of you
loecherPath = "Teaching/SS2023/IntroCoding/datasets/"
df = pd.read_csv(rootPath + loecherPath + "train.csv")
#df = pd.read_csv('/content/drive/MyDrive/Teaching/SS2023/IntroCoding/datasets/train.csv')
```

```
#or
url = "https://drive.google.com/file/d/1hzvcubf2B8PKtjG40AcytQKw0lESkBVW/view?usp=sharing"
url='https://drive.google.com/uc?id=' + url.split('/')[2]
df = pd.read_csv(url)
df.head()
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
#   Column          Non-Null Count  Dtype
#   ...
```

0	Id	1460 non-null	int64
1	MSSubClass	1460 non-null	int64
2	MSZoning	1460 non-null	object
3	LotFrontage	1201 non-null	float64
4	LotArea	1460 non-null	int64
5	Street	1460 non-null	object
6	Alley	91 non-null	object
7	LotShape	1460 non-null	object
8	LandContour	1460 non-null	object
9	Utilities	1460 non-null	object
10	LotConfig	1460 non-null	object
11	LandSlope	1460 non-null	object
12	Neighborhood	1460 non-null	object
13	Condition1	1460 non-null	object
14	Condition2	1460 non-null	object
15	BldgType	1460 non-null	object
16	HouseStyle	1460 non-null	object
17	OverallQual	1460 non-null	int64
18	OverallCond	1460 non-null	int64
19	YearBuilt	1460 non-null	int64
20	YearRemodAdd	1460 non-null	int64
21	RoofStyle	1460 non-null	object
22	RoofMatl	1460 non-null	object
23	Exterior1st	1460 non-null	object
24	Exterior2nd	1460 non-null	object
25	MasVnrType	1452 non-null	object
26	MasVnrArea	1452 non-null	float64
27	ExterQual	1460 non-null	object
28	ExterCond	1460 non-null	object
29	Foundation	1460 non-null	object
30	BsmtQual	1423 non-null	object
31	BsmtCond	1423 non-null	object
32	BsmtExposure	1422 non-null	object
33	BsmtFinType1	1423 non-null	object
34	BsmtFinSF1	1460 non-null	int64
35	BsmtFinType2	1422 non-null	object
36	BsmtFinSF2	1460 non-null	int64
37	BsmtUnfSF	1460 non-null	int64
38	TotalBsmtSF	1460 non-null	int64
39	Heating	1460 non-null	object
40	HeatingQC	1460 non-null	object
41	CentralAir	1460 non-null	object

42	Electrical	1459	non-null	object
43	1stFlrSF	1460	non-null	int64
44	2ndFlrSF	1460	non-null	int64
45	LowQualFinSF	1460	non-null	int64
46	GrLivArea	1460	non-null	int64
47	BsmtFullBath	1460	non-null	int64
48	BsmtHalfBath	1460	non-null	int64
49	FullBath	1460	non-null	int64
50	HalfBath	1460	non-null	int64
51	BedroomAbvGr	1460	non-null	int64
52	KitchenAbvGr	1460	non-null	int64
53	KitchenQual	1460	non-null	object
54	TotRmsAbvGrd	1460	non-null	int64
55	Functional	1460	non-null	object
56	Fireplaces	1460	non-null	int64
57	FireplaceQu	770	non-null	object
58	GarageType	1379	non-null	object
59	GarageYrBlt	1379	non-null	float64
60	GarageFinish	1379	non-null	object
61	GarageCars	1460	non-null	int64
62	GarageArea	1460	non-null	int64
63	GarageQual	1379	non-null	object
64	GarageCond	1379	non-null	object
65	PavedDrive	1460	non-null	object
66	WoodDeckSF	1460	non-null	int64
67	OpenPorchSF	1460	non-null	int64
68	EnclosedPorch	1460	non-null	int64
69	3SsnPorch	1460	non-null	int64
70	ScreenPorch	1460	non-null	int64
71	PoolArea	1460	non-null	int64
72	PoolQC	7	non-null	object
73	Fence	281	non-null	object
74	MiscFeature	54	non-null	object
75	MiscVal	1460	non-null	int64
76	MoSold	1460	non-null	int64
77	YrSold	1460	non-null	int64
78	SaleType	1460	non-null	object
79	SaleCondition	1460	non-null	object
80	SalePrice	1460	non-null	int64

dtypes: float64(3), int64(35), object(43)

memory usage: 924.0+ KB

```
df.head()
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPu
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPu
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPu
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPu
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPu

6.0.1 Tasks:

1. Identify the columns with missing values (Hint: use the `any` function) and read up their description on the [kaggle](#) site
2. Replace missing values with “appropriate” values, as follows:
 - for “categorical” data (e.g. strings) use the most frequent value (`mode`)
 - for numerical data: plot a histogram and look at the distribution. For rather symmetric looking data, choose the mean, otherwise the median.
 - for “time” variables such as year: find another year variable as a proxy (Hint: read up on the [combine_first](#) function)
3. Find those columns with fewer than 8 unique values (Hint: use the pandas method `nunique()`)
 - Create 2 insightful boxplots: SalePrice versus YrSold or MSZoning. Decide if a log scale would be more discerning.
 - Use `groupby` to compute the boxes, i.e. the lower and upper quartiles. (Hint: use the numpy or pandas method `quantile`)
 - Then compute the whiskers
 - And find the outliers

7 Regression, seaborn

Correlation and Regression as well as a quick exploration of the seaborn visualization capabilities

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from numpy.random import default_rng
#!pip install gapminder
#from gapminder import gapminder

#new library
import statsmodels.api as sm
import statsmodels.formula.api as smf
```

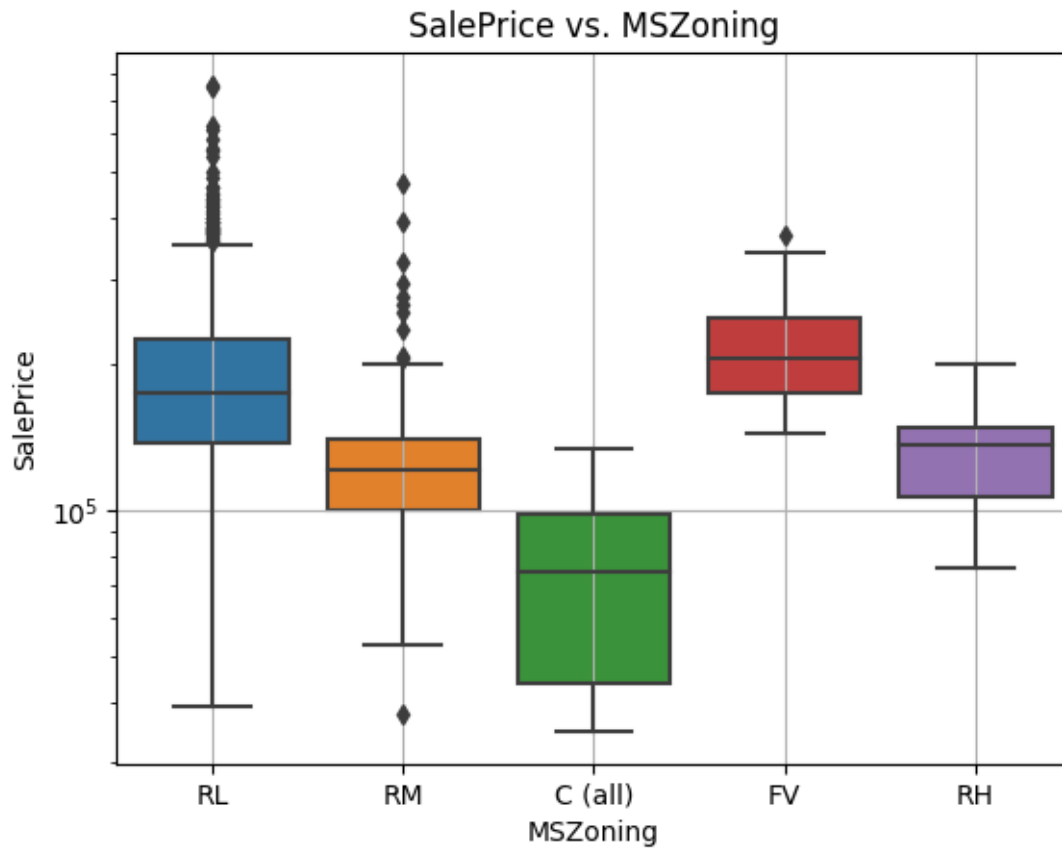
[kaggle Housing Data](#)

```
#or
url = "https://drive.google.com/file/d/1hzvcubf2B8PKtjG40AcytQKw0lESkBVW/view?usp=sharing"
url = 'https://drive.google.com/uc?id=' + url.split('/')[2]
df = pd.read_csv(url)
df.head()
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPu
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPu
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPu
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPu
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPu


7.1 Seaborn Graphs

```
sns.boxplot(df, y = "SalePrice", x = "MSZoning");  
plt.yscale("log");plt.grid();  
plt.title("SalePrice vs. MSZoning");
```




7.1.0.1 Multiple Groups

```
sns.boxplot(df, y = "SalePrice", x = "MSZoning", hue = "YrSold");  
plt.yscale("log");  
plt.grid();  
plt.title("SalePrice vs. MSZoning");
```




7.2 Violin Plots

```
sns.violinplot(df, y = "SalePrice", x = "MSZoning");
plt.yscale("log");plt.grid();
plt.title("SalePrice vs. MSZoning");
```




7.2.1 Histograms

```
sns.histplot(data=df, x="SalePrice", kde=True);plt.grid();
```



```
sns.histplot(data=df, x="SalePrice", kde=True, hue = "MSZoning");plt.grid();
```



7.2.2 Tasks:

7.2.2.1 Regression/Correlation (Housing Data)

1. Look up the `pairplot` function and create pairwise scatter plots of
 - 5-7 hand-picked numerical features, one of them being `SalePrice`
 - Hint: look at `dtypes`
2. Choose the row with `SalePrice` and pick two reasonably strong correlations.
 - Compute the correlation coefficients
 - Fit a simple regression line (with `statsmodels`) for each and visualize them using `regplot`
 - Fit a **multiple regression** by including both *explanatory variables* and compare the coefficients

```
df.dtypes != "object"
```

```

Id                True
MSSubClass        True
MSZoning          False
LotFrontage       True
LotArea           True
...
MoSold            True
YrSold            True
SaleType          False
SaleCondition     False
SalePrice         True
Length: 81, dtype: bool

```

```
df.columns[df.dtypes != "object"][1:]
```

```

Index(['MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual', 'OverallCond',
      'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2',
      'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
      'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
      'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces',
      'GarageYrBlt', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF',
      'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal',
      'MoSold', 'YrSold', 'SalePrice'],
      dtype='object')

```

7.3 Extra Credit

7.3.0.1 Modeling Missing Values Titanic Data

1. detect the missing values
2. replace the NAs in survived with the estimate grouped by sex

```

#titanic
titanic = sns.load_dataset('titanic')
titanic["3rdClass"] = titanic["pclass"]==3
titanic["male"] = titanic["sex"]=="male"
#titanic.head()

#Introduce some missing values
rng = default_rng()

```

```
missingRows = rng.integers(0,890,20)
print(missingRows)
#introduce missing values
titanic.iloc[missingRows] = np.nan
```

8 String Exercises

(Lab 8)

8.1 Part I

1. Write a function that takes a string as input and returns the string with all the vowels removed.
2. Write a function that takes a string as input and returns the number of words in the string. Assume that words are separated by spaces.
3. Write a function that takes a string as input and returns a new string with the words reversed. Assume that words are separated by spaces.
4. Write a function that takes a string as input and checks if it is a palindrome (reads the same forwards and backwards).

8.1.1 NOTE!

- The skeleton code is already prepared for you below
- Do **not** just blindly copy code from the internet.
- Use your new skills, not some specialized functions which you do not understand

8.2 Part II (Netflix data)

1. How many unique shows are there for user Olivia ?

From here on it is all about Olivia only.

2. What is the most frequent show ?
3. Find the show with the shortest string length. (Hint: Use list comprehension)
4. Find the show with the **most** number of words. (Hint: Use list comprehension)
5. How many show names contain a number ?
6. How many show names end in a number at the end ?

7. Extract those titles which contain the pattern (Episode _) where _ is a number.

- Split those strings on the : and extract only the first part
- Remove duplicates
- How many of those strings are also in the Showcolumn?

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import re
```

```
x = [1,2,2,3,3,3,3,3,3]
vals, cts = np.unique(x, return_counts=True)
vals[cts == np.max(cts)]
```

```
array([3])
```

```
# Exercise 1:
# Write a function that takes a string as input and returns the string with all the vowels removed
#Hint: removing strictly speaking is a replace
```

```
def remove_vowels(string):
```

```
# Example usage
text = "Hello, World!"
result = remove_vowels(text)
print(result) # Output: Hll, Wrld!
```

```
re.sub("\d", "", "My 1st password is 1234")
```

```
'My st password is '
```

```
re.sub("[a-l]", "", "My 1st password is 1234") # "" replaces with "nothing"
```

```
'My 1st psswor s 1234'
```

```
re.sub("[aeiou]", "", "My 1st password is 1234")
```

```
def remove_vowels(s):
```

```
print(remove_vowels("Hello, World!"))  
print(remove_vowels("Berlin School of Economics"))
```

Hll, Wrld!

Brln Schlf cmcs

```
len("Berlin School of Economics".split(sep= " "))#this one splits on exactly one white space  
re.split(r"\s+", "Berlin School of Economics")#this one splits on exactly one or more whitespaces
```

```
['Berlin', 'School', 'of', 'Economics']
```

```
#Exercise 2:
```

```
# Write a function that takes a string as input and returns the number of words in the string
```

```
def count_words(string):
```

```
# Example usage
```

```
text = "This is a sample sentence."
```

```
result = count_words(text)
```

```
print(result) # Output: 5
```

```
#Exercise 3:
```

```
# Write a function that takes a string as input and returns a new string with the
```

```
# characters of the words reversed. Assume that words are separated by spaces.
```

```
def reverse_words(string):
```

```
# Example usage
```

```
text = "Hello, World!"
```

```
result = reverse_words(text)
```

```
print(result) # Output: olleH, dlroW!
```

```
#Exercise 4:
```

```
# Write a function that takes a string as input and checks if it is a palindrome (reads the same forwards and backwards)
```

```
def is_palindrome(string):  
  
# Example usage  
text = "radar"  
result = is_palindrome(text)  
print(result) # Output: True
```

8.3 Netflix

```
url = "https://drive.google.com/file/d/1-1rqPVMKh3LMviUGyG1MfQpFTtn0rJ5V/view?usp=sharing"  
url='https://drive.google.com/uc?id=' + url.split('/')[2]  
#reading the zipped version seems not to be working  
netflix = pd.read_csv(url)#, compression="zip")
```

9 Subplots and Pivoting

(Lab 9)

1. Smoking Data

- Group by age (rounded to 10 years) **and** smoking status and compute the survival rates separately for the two groups.
- Compute the average difference in survival.
- Create a grid of subplots with 1 row and 2 columns.
 - Left plot: Line plot of `p_alive` and `p_smokes` as a function of age. Annotate the maximum `p_smokes`.
 - Right plot: Line plot of survival rates for the two smoking groups. Annotate the maximum difference.

2. Netflix Data

- Create a grid of subplots with 2 rows and 1 column.
 - Top plot: Create a **daily** plot of hours spent watching ‘Modern Family’
 - Extra Credit: Find the top peak and annotate it by the user watching most that day (plus his/her total duration).
 - Bottom plot: Create a **weekly** plot of hours spent watching ‘Modern Family’ (Hint: recall the `resample` command)

If you need a refresher on `resample`, here is the Data Camp course on [Working with dates and times in python](#)

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from datetime import datetime
```

9.1 Smoking Data

3 columns: “outcome” measures whether the person is still alive after 10 years.

We want to glean the effect of smoking on survival probability.

```
df = pd.read_csv("https://calmcode.io/datasets/smoking.csv")
df = df.assign(alive = (df['outcome'] == 'Alive').astype(int),
               smokes = (df['smoker'] == 'Yes').astype(int))
df.head()
```

	outcome	smoker	age	alive	smokes
0	Alive	Yes	23	1	1
1	Alive	Yes	18	1	1
2	Dead	Yes	71	0	1
3	Alive	No	67	1	0
4	Alive	No	64	1	0

Smoking is good for your health ??

```
df.groupby(['smoker']).alive.mean()
```

```
smoker
No      0.685792
Yes     0.761168
Name: alive, dtype: float64
```

```
df.groupby(['smoker']).agg(prob=('alive', np.mean))
```

	prob
smoker	
No	0.685792
Yes	0.761168

Age adjustment

```
np.round(19 / 10) * 10
```

20.0

```
df = df.assign(age10 =np.round(df['age'] / 10) * 10)
df
```

	outcome	smoker	age	alive	smokes	age10
0	Alive	Yes	23	1	1	20.0
1	Alive	Yes	18	1	1	20.0
2	Dead	Yes	71	0	1	70.0
3	Alive	No	67	1	0	70.0
4	Alive	No	64	1	0	60.0
...
1309	Alive	Yes	35	1	1	40.0
1310	Alive	No	33	1	0	30.0
1311	Alive	Yes	21	1	1	20.0
1312	Alive	No	46	1	0	50.0
1313	Alive	Yes	41	1	1	40.0

```
df2 = (df.groupby(['age10'])
       .agg(p_alive=('alive', np.mean)
            ,p_smokes=('smokes', np.mean)))
df2
```

	p_alive	p_smokes
age10		
20.0	0.980645	0.445161
30.0	0.972332	0.438735
40.0	0.900000	0.495833
50.0	0.821622	0.632432
60.0	0.587302	0.464286
70.0	0.203947	0.236842
80.0	0.000000	0.168831

9.2 Netflix Data

```
url = "https://drive.google.com/uc?id=1-1rqPVMKh3LMviUGyG1MfQpFTtn0rJ5V"
netflix = pd.read_csv(url, parse_dates=["Start.Time", "Start.Day"])
#netflix['Start.Day'] = pd.to_datetime(netflix['Start.Day'])
netflix.head()
```

	Profile	Start.Time	Duration	Attributes	Title
0	Olivia	2021-01-17 16:22:20	32	NaN	Ginny & Georgia: Season 2: Latkes Are Lit (Epi...
1	Olivia	2021-01-17 16:07:48	126	NaN	Ginny & Georgia: Season 2: Latkes Are Lit (Epi...
2	Olivia	2021-01-17 15:55:31	385	NaN	Ginny & Georgia: Season 2: Latkes Are Lit (Epi...
3	Olivia	2021-01-17 13:02:24	3426	NaN	Ginny & Georgia: Season 2: Happy My Birthday t.
4	Olivia	2021-01-17 12:08:05	3257	NaN	Ginny & Georgia: Season 2: What Are You Playin.